

# CPSC 424/524 Assignment 1

---

yl2335, Yangxiaokang Liu

## Environment

```
[cp424_yl2335@grace2 ~]$ module load intel/2018b
[cp424_yl2335@grace2 ~]$ module list

Currently Loaded Modules:
  1) StdEnv                (S)  3) binutils/2.30-GCCcore-7.3.0    5) ifort/2018.3.222-GCC-7.3.0-2.30
  7) impi/2018.3.222-iccifort-2018.3.222-GCC-7.3.0-2.30  9) imkl/2018.3.222-iimpi-2018b
  2) GCCcore/7.3.0         4) icc/2018.3.222-GCC-7.3.0-2.30  6) iccifort/2018.3.222-GCC-7.3.0-2.30
  8) iimpi/2018b          10) intel/2018b

Where:
  S: Module is Sticky, requires --force to unload or purge

[cp424_yl2335@grace2 ~]$ which icc
/gpfs/loomis/apps/avx/software/ifort/2018.3.222-GCC-7.3.0-2.30/compilers_and_libraries_2018.3.222/linux/bin/intel64/icc
[cp424_yl2335@grace2 ~]$ icc --version
icc (ICC) 18.0.3 20180410
Copyright (C) 1985-2018 Intel Corporation. All rights reserved.
```

## Running Experiments

### Command

```
srun --pty -p interactive -c 1 -t 6:00:00 --mem-per-cpu=6gb bash ./run.sh
```

### Outputs

-----Make-----

rm ex1a ex1b ex1c ex2

icc -g -O0 -fno-alias -std=c99 -o ex1a ex1.c timing.o

icc -g -O3 -xHost -fno-alias -std=c99 -o ex2 ex2.c dummy.c timing.o

icc -g -O3 -no-vec -no-simd -fno-alias -std=c99 -o ex1b ex1.c timing.o

icc -g -O3 -xHost -fno-alias -std=c99 -o ex1c ex1.c timing.o

-----Ex1 Option a-----

Approximation = 3.141593

Sin(4 \* area) = 0.000000

Millions of flops: 5000.000000

Time consumed: 5.434563

Millions of flops per second: 920.037150

-----Ex1 Option b-----

Approximation = 3.141593

Sin(4 \* area) = 0.000000

Millions of flops: 5000.000000

Time consumed: 5.434214

Millions of flops per second: 920.096285

-----Ex1 Option c-----

Approximation = 3.141593

Sin(4 \* area) = -0.000000

Millions of flops: 5000.000000

Time consumed: 2.751773

Millions of flops per second: 1817.010264

-----Ex2-----

N Time Mflops

9 1.772019 1363.370950

19 1.670107 3053.860625

40 1.350091 3976.553579

85 1.205822 4730.593307

180 1.130042 5344.754755

378 1.148801 5520.353099

```

794 1.210063  5504.304204
1667  1.171854  2983.266112
3502  1.235782  2971.489652
7355  1.465514  2631.253189
15447 1.325071  1527.970227
32439 1.414168  1503.302383
68122 1.483196  1505.007879
143056 1.573838  1489.244453
300419 1.647766  1493.557116
630880 1.038121  1244.596977
1324849 1.028950  659.237770
2782184 1.058159  673.092785
5842587 1.134145  659.396393
12269432 1.239925  633.299358
25765808 1.205754  683.809324
54108198 1.355506  638.677515
113627216 1.392969  652.575756

```

## Exercise 1: Division Performance


1. Demonstrate that your program actually computes a reasonable approximation to  $\pi$  .
  - a. The outputs printed as doubles are 3.141593 which is close to the true value of  $\pi$ .
  - b.  $\sin(4 * area) = 0$  which means approximation must be multiples of  $\pi$ . We already know from (a) that it can't be other multiples but just  $\pi$ .
2. Report performances of option (a), option (b) and option (c) and explain their differences.
  - Phenomena
 

The MFlops of option (a) and option (b) are almost the same. While the MFlops of option (c) almost doubles.
  - Explanation

Both option (b) and option (c) uses -O3. According to the manual icc, **O3 performs O2 optimizations and enables more aggressive loop transformations such as Fusion, Block-Unroll-and-Jam, and collapsing IF statements. However, option (b) uses -no-vec which disables all auto-vectorization, including vectorization of array notation statements. The option -no-simd disables vectorization of loops that have simd pragmas.** These two options make option (b) as slow as option (a). Option (c) uses -xHost which tells the compiler to generate instructions for the highest instruction set available on the compilation host processor. "Ivy Bridge" nodes use AVX (Advanced Vector Extensions) instruction set which widen the vector registers from 128 bits to 256 bits, so the floating-point hardware can sustain 16 single-precision or 8 double-precision floating point operations per cycle, MFlops therefore doubles.

## Exercise 2: Vector Triad Performance

1. Generate a plot of the performance in MFlops vs. N. Use a logarithmic scale for N on the x-axis.




2. Explain the variations in terms of the processor architecture, by computing and discussing the apparent memory bandwidth that your data show near the variations in your plot.

a. Use `lscpu` to list out the information of the cpu

*Architecture:* x86\_64  
*CPU op-mode(s):* 32-bit, 64-bit  
*Byte Order:* Little Endian  
*CPU(s):* 16  
*On-line CPU(s) list:* 0-15  
*Thread(s) per core:* 2  
*Core(s) per socket:* 8  
*Socket(s):* 1  
*NUMA node(s):* 1  
*Vendor ID:* GenuineIntel  
*CPU family:* 6  
*Model:* 62  
*Model name:* Intel(R) Xeon(R) CPU E5-2650 v2 @ 2.60GHz  
*Stepping:* 4  
*CPU MHz:* 2600.016  
*BogoMIPS:* 5200.03  
*Virtualization:* VT-x  
*L1d cache:* 32K  
*L1i cache:* 32K  
*L2 cache:* 256K  
*L3 cache:* 20480K  
*NUMA node0 CPU(s):* 0-15

b. We can also draw the performane in Bandwidth vs. N.



- c. From the data we can identify three major drops of the bandwidth
- i. Bandwidth dropped from 88 GB/s to 48 GB/s when N increases from 794 to 1667. **The bandwidth drops because when  $N = 794$ , the size of the vectors is  $4 * N * \text{sizeof}(\text{double}) = 25408$  Bytes which is smaller than the capacity of the L1 cache 32 KB, when  $N = 1667$ , the size of the vectors is 53344 Bytes which exceeds the capacity of the L1d cache. Cache misses cause the bandwidth to drop.**
  - ii. Bandwidth dropped from 42 GB/s to 24 GB/s when N increases from 7335 to 15447. **When  $N = 7335$ , the size of the vectors is 234720 B which is smaller than the capacity of L2 cache 256KB, when  $N = 15447$ , the size of the vectors is 494304 B which exceeds L2's capacity. The bandwidth dropped.**
  - iii. Bandwidth dropped from 20 GB/s to 11 GB/s when N increases from 630880 to 1324849. **When  $N = 630880$ , the size of the vectors is 19715.0 KB which is smaller than the capacity of L3 cache 20480 KB, when  $N = 1324849$ , the size of the vectors is 41401.53125 KB which exceeds L3's capacity. The bandwidth dropped.**